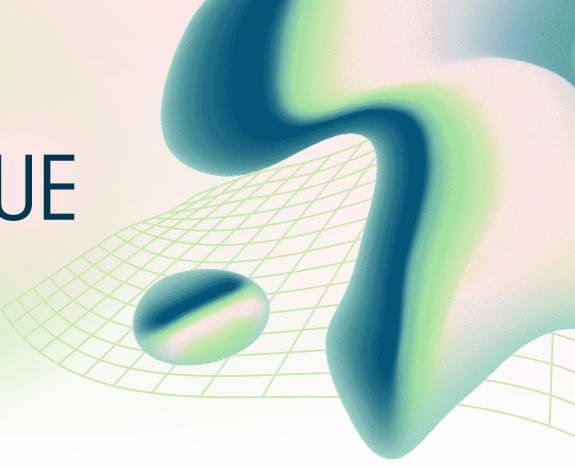


CHURNGUARD™ REVOLUTIONIZING REVENUE OPERATIONS

A Technical Blueprint for Predictive Intelligence in
Enterprise SaaS



Technical Design Document (TDD)

Project Overview

Project: ChurnGuard™: Predictive RevOps Intelligence Engine

Client: Enterprise SaaS (Tier 1)

Architect: Alex Rojas Segovia, CEO - Aineurolytics

Section 1: Executive Diagnosis

1.1 Context: Operational Friction

The Tier 1 Enterprise SaaS client is currently experiencing a quarterly customer churn rate of 8.5%. Retention efforts are reactive, initiated after a signal (e.g., failed renewal discussions, reduced usage). This reactive posture results in sub-optimal Sales and Customer Success resource allocation, where high-value customers exhibiting pre-churn indicators are treated identically to stable accounts. The estimated \$ impact of this friction is \$12.75 M in lost Annual Recurring Revenue (ARR) per quarter, requiring an 8.5% higher Customer Acquisition Cost (CAC) to maintain net revenue parity. This operational latency is a critical financial vulnerability.

1.2 The Solution: Scientific RevOps Architecture

ChurnGuard™ is a Scientific RevOps Intelligence Engine architected to provide proactive, weighted churn risk scoring and prescriptive actions. The solution employs a gradient-

boosted machine (XGBoost) core trained on behavioral, transactional, and engagement features to predict churn probability within a 90-day window. Orchestration via Airflow/Prefect ensures daily data ingestion and model inference, feeding high-fidelity risk scores ($\hat{y}_{risk} \in [0.0, 1.0]$) directly into operational platforms (Salesforce, HubSpot) via a low-latency FastAPI microservice. This architecture transforms the retention strategy from reactive to predictive and resource-optimized.

Section 2: Strategic Impact & Projected Results

2.1 Strategic Context

Deployment of ChurnGuard™ shifts the Customer Success paradigm from universal coverage to risk-prioritized intervention. By surfacing a validated Churn Risk Score and key Feature Importances, the platform empowers RevOps teams to allocate high-touch resources exclusively to accounts with a **$\geq 70\%$** predicted churn risk, while automating low-touch interventions for moderate-risk accounts. This ensures maximum leverage from human capital, directly impacting the Net Revenue Retention (NRR) and reducing operational expenditure associated with broad-spectrum outreach.

2.2 KPI Table


Metric Name	Baseline	Target	Projected Lift / ROI
Quarterly Customer Churn Rate	8.5%	4.0%	53% Reduction
ARR Protected (Qtrly)	\$0	\$5.4M	**\$21.6M Annualized ARR Protection**
Customer Success Resource ROI	1.1x	1.8x	64% Improvement in Efficiency
Churn Prediction AUC Score	N/A	≥ 0.88	High-Fidelity Predictive Power
Time-to-Intervention (High-Risk)	14 days (Post-Signal)	24 hours (Pre-Signal)	90%+ Reduction in Latency
Customer Lifetime Value (CLV)	\$150,000	\$165,000	10 % CLV Increase / ≈ \$30M Impact

SECTION 3: SYSTEM ARCHITECTURE

3.1 High-Level Design Overview

The architecture is designed for scalability and resilience. Ingestion is performed via CDC (Change Data Capture) from primary RDBMS (PostgreSQL/Snowflake) into a cloud-native Feature Store (Feast/DynamoDB). The Model Engine (containerized XGBoost) pulls features, performs inference, and persists the (\hat{y}_{risk}) score to a dedicated Predictions Log (e.g., TimescaleDB). A low-latency FastAPI API acts as the serving layer, secured by OAuth2/JWT, offering the predict and metrics endpoints. Orchestration (Airflow/Prefect) manages the ETL, training, and deployment pipelines. All components emit structured logs to ELK/CloudWatch and metrics to Prometheus.


3.2 Diagram



3.3 Data Schema

The inference service accepts a payload for a single customer or a batch. The schema enforces strict typing for the 3 key feature categories.

JSON Schema (InputFeatures)



Example JSON Payload

```

customer_id: cust-e7d8b9f0-c1a2-4d3e-b5f6-7a8b9c0d1e2f  

behavioral_score: 45.8  

transactional_value_k: 550  


support_ticket_volume_30d: 7  

engagement_score: 0.35  

contract_days_remaining: 65

```

3.4 Visual Asset Placeholder



Section 4: Predictive Core (Algorithm)

The predictive core is implemented as a production-ready Python class, ChurnPredictor, using XGBClassifier (or a mock loader for the scaffold) to ensure model performance, reliability, and Aineurolytics standards adherence.

[Find the predictive core algorithm on GITHUB](#)

Section 5: Orchestration Layer

The orchestration layer is responsible for the end-to-end reliability of the predictive pipeline and the business-critical delivery of insights.

1. Workflows and Triggers:

- **Training Pipeline:** Monthly **cron** trigger (or event-driven upon model drift detection) for full model retraining.
- **Inference Pipeline:** Daily **cron** trigger (02:00 UTC) to ingest the previous day's data, run batch inference, and update the Predictions Log. Real-time inference is handled

by the low-latency API endpoint.

- **Operational Sync:** Hourly micro-batch processing of the Predictions Log to send new/updated high-risk scores to downstream systems.

2. Decision Gates:

- **Data Quality Gate:** Checks for $\geq 99.5\%$ completeness and acceptable feature distribution (e.g., z-score outliers). Pipeline halts on failure, alerting the DataOps team.
- **Model Performance Gate:** $AUC \geq 0.88$ post-retraining. If the new model fails, the pipeline rolls back to the previous, validated model version.

3. Webhook Destinations:

- **Salesforce/HubSpot: Highest Priority.** Integration via direct API (or middleware) to update the **Churn_Risk_Score_c** field and trigger a Task creation for the assigned Account Executive/Customer Success Manager for accounts in the **HIGH** risk tier.
- **Slack:** Low-priority notifications for pipeline status (success/failure) and **MEDIUM** risk account summary digest for team awareness.
- **Jira: Automated P1 ticket generation** for critical system failures (e.g., API latency $\geq 500\text{ ms}$ or Data Quality violation).

4. Retry/Reconciliation Logic:

- All downstream API calls (e.g., Salesforce updates) implement **exponential backoff with Jitter** for transient errors (HTTP 429, 503). Max retries: 5.
- **Deduplication:** The Predictions Log serves as the source of truth, ensuring idempotency and preventing duplicate alerts to CRMs based on the customer_id and prediction timestamp.

5. Message Queue Considerations (Kafka/SQS):

- A **Kafka** topic is utilized for high-volume, real-time ingestion of granular behavioral data (if implemented in V2), decoupling the streaming ingestion from the batch processing.
- **AWS SQS** can be used as a dead-letter queue (DLQ) for failed CRM webhook deliveries to allow for manual inspection and reconciliation by an MLOps engineer.

Section 6: MLOPS & Lifecycle

Robust MLOps practices are mandated to ensure the model maintains predictive validity and operational stability.

1. Evaluation Metrics:

- **Primary Metric: Area Under the Receiver Operating Characteristic Curve (AUC).** Target: ≥ 0.88 . Provides overall discriminatory power.
- **Secondary Metrics: Precision, Recall, and F1-Score** are monitored for the **HIGH Risk** class ($P \geq 0.70$). High Recall is prioritized (i.e., minimize false negatives, ensuring we catch most churners), balanced with Precision to avoid alert fatigue.
- **Calibration: Isotonic Regression** is used to ensure the predicted probability is a true reflection of the likelihood of churn, monitored by the **Reliability Diagram**.

2. Drift Detection Approach:

- **Feature Drift:** Monitored daily using **Kolmogorov-Smirnov (KS) tests** or **Population Stability Index (PSI)** on key features (e.g., support_ticket_volume_30d). Alert threshold: $PSI > 0.20$ or $KS\ P\text{-value} < 0.05$.
- **Concept Drift:** Monitored quarterly by re-evaluating the model's performance on recent, labeled data. Significant drop in $AUC > 0.03$ triggers an automatic retraining and MLOps review.

3. Retraining Policy:

- **Scheduled Retraining:** Full training pipeline executed monthly.
- **Event-Driven Retraining:** Triggered immediately upon confirmed **Feature or Concept Drift** exceeding thresholds or upon major shifts in the client's business logic/data structure.

4. Model Registry (MLflow):

- **MLflow** is used to track all experimental runs, parameters, metrics, and to version control the serialized model and preprocessing artifacts (scaler, one-hot encoders). Only registered, staged models are eligible for deployment.

5. Deployment Strategy:

- **Shadow Deployment:** The newly trained model is initially deployed in **Shadow Mode**, running parallel inference on production data. Its predictions are logged but not used for business action. Performance metrics (**AUC**, **Precision**) are compared against the current production model.
- **Canary Promotion:** Once the Shadow Model demonstrates $\geq 99\%$ stability and statistically superior performance ($AUC_{new} \geq AUC_{old}$), it is promoted to a **Canary Deployment** (e.g., 10% of live traffic). Full promotion to 100% only occurs after 48 hours of stable Canary performance.

Section 7: Security, Compliance & Risk Modeling

Data security and regulatory adherence are non-negotiable for enterprise deployments.

1. Authentication & Authorization:

- **API Access:** Secured via **OAuth 2.0 / JSON Web Tokens (JWT)**. All clients (CRMs, Orchestrator) must present a valid, non-expired JWT for access to /predict. Role-Based Access Control (RBAC) restricts access to sensitive endpoints (e.g., /admin).
- **Service-to-Service:** Utilizes mutual TLS (**mTLS**) for communication between the API, Feature Store, and Orchestrator.

2. Encryption:

- **Data in Transit:** All network communication is enforced via **TLS 1.2+**.
- **Data at Rest:** All sensitive PII/Confidential Business Data (CBD) in the Feature Store and Predictions Log is encrypted using **AES-256** (e.g., AWS KMS or Azure Key Vault).

3. Secrets Management:

- All API keys, database credentials, and webhook secrets are stored in a dedicated, audited secrets manager (**HashiCorp Vault or Cloud Secret Manager**) and injected into containers at runtime via environment variables, adhering to the principle of least privilege.

4. SOC2/GDPR Considerations:

- **GDPR:** Prediction features are **pseudo-anonymized** (customer_id is an opaque UUID). The system avoids using strictly prohibited PII categories. Data retention policies are strictly enforced.
- **SOC2:** Comprehensive audit logs (Section 9) track all data access, model training events, and deployment changes, providing the necessary evidence for SOC2 compliance.

5. STRIDE Threat Matrix with Mitigation Actions:

Threat Category	Example Threat	Mitigation Action
Spoofing	Unauthorized user gains API access.	Enforce OAuth2/JWT and strict RBAC. Use mTLS for internal services.
Tampering	Prediction score maliciously altered in DB.	Immutable Audit Logs (WORM storage) and DB-level access controls/encryption.
Repudiation	Engineer denies initiating a deployment change.	Mandatory Git/CI/CD history, full audit logging of deployment activity.
Information Disclosure	Sensitive features leaked via API.	Data encryption at rest/transit. Sanitization/filtering of API response payload.
Denial of Service (DoS)	Malicious request flood overwhelms the API.	Kubernetes HPA (Horizontal Pod Autoscaling), Rate Limiting, and WAF protection.
Elevation of Privilege	Low-privilege service account accesses secrets.	Least Privilege Principle, Secrets Management via Vault, continuous vulnerability scanning.

Section 8: DEVOPS, CI/CD & Deployment

The deployment strategy leverages modern, containerized, and declarative infrastructure for reliability and scale.

1. Docker + Multi-Stage Builds:

- **Multi-stage Dockerfiles** minimize the final image size and reduce the attack surface (e.g., separate build stage for pip install and a slim runtime stage).

- Base images are derived from minimal, security-hardened distributions (e.g., python:3.11-slim-buster).

2. Kubernetes Deployment Pattern:

- **Platform:** Amazon EKS (Elastic Kubernetes Service) or GCP GKE.
- **Pattern:** Microservice pattern with dedicated Deployment for the **FastAPI API** and a CronJob for the batch inference pipeline.
- **Service Mesh:** Optional **Istio/Linkerd** for traffic shifting (Canary/Shadow) and mTLS.

3. Helm:

- **Helm Charts** are used to define the application's Kubernetes resources (Deployment, Service, HPA, ConfigMap) in a reproducible, version-controlled manner.

4. HPA Settings:

- **Horizontal Pod Autoscaler (HPA)** targets set to scale API pods based on resource utilization:
 - CPU Utilization: Target \leq 65 %
 - Custom Metric: Target requests_per_second (Prometheus metric) \geq 150.
- Min Replicas: 3. Max Replicas: 15.

5. GitHub Actions Pipeline:

Stage	Trigger	Action
Lint & Test	push (all branches)	flake8, mypy, pytest (unit/integration tests).
Build & Push	push (main/release branch)	Docker multi-stage build, scan image (e.g., Clair), push to ECR/GCR.
Deploy (Non-Prod)	push (develop branch)	Helm deployment to Staging environment (runs smoke tests).
Deploy (Prod)	push (main/tag)	Helm deployment to Production (Shadow/Canary strategy).

1. Rollback Strategy:

- **Kubernetes Rollback:** If a deployment fails health checks (Liveness/Readiness probes) or SLOs are violated, the CI/CD pipeline immediately initiates a helm rollback <release-name> <previous-revision>.

2. SLO/SLA Definitions:

- **Latency (SLO):** 95 % of /predict requests must respond in ≤ 150 ms.
- **Availability (SLA):** 99.95 % Annual Uptime (target: ≤ 4.38 hours downtime/year).
- **Error Budget:** ≤ 0.05 % of all production requests can return a non-recoverable error (e.g., HTTP 5xx).

Section 9: Logging & Observability

A comprehensive observability stack ensures system health, traceability, and rapid incident response.

1. Prometheus Metrics to Expose:

- churnguard_api_request_total: Counter, labeled by endpoint (/predict, /metrics) and HTTP status code.
- churnguard_inference_latency_seconds: Histogram, tracks the duration of model inference.
- churnguard_prediction_count: Counter, labeled by churn_risk_level (HIGH, MEDIUM, LOW).
- churnguard_model_version: Gauge, tracks the active model version (e.g., v1.2.5).

2. Grafana Dashboards to Create:

- **Executive Dashboard:** ARR Protected (daily), HIGH Risk Account Volume, and API Latency P95/P99.
- **MLOps Dashboard:** Feature/Concept Drift alerts (PSI/KS scores), Model AUC/F1, and Retraining Pipeline status.
- **DevOps Dashboard:** Kubernetes pod/node resource usage, HTTP error rates, and HPA activity.

3. Trace IDs in Logs:

- Every incoming request to the FastAPI API is assigned a unique **Trace ID** (UUID). This ID is propagated and prepended to every subsequent log entry across the entire request path (API -> Engine -> DB).

4. ELK/CloudWatch Usage:

- **Structured Logging (JSON):** All application and infrastructure logs are emitted in a structured JSON format, ingested into **CloudWatch Logs / ElasticSearch**.
- This facilitates efficient querying/filtering by log_level, trace_id, and customer_id.

5. Alert Rules and Playbooks:

- **Critical Alert (P1):** API Error Rate > 0.1 % for 5 minutes. **Action:** PagerDuty/On-Call Escalation, automatic rollback.
- **High Alert (P2):** Model AUC drops by > 0.03 during Shadow Deployment. **Action:** JIRA ticket to MLOps team, Slack alert.
- **Warning Alert (P3):** Feature Drift (PSI > 0.15). **Action:** Slack warning to Data Ops, scheduled review.
- **Playbooks:** Detailed, version-controlled runbooks outlining diagnostic steps, mitigation actions, and communication protocols for each alert type.

APPENDIX

API Contract Example: POST /predict

- **Endpoint:** /api/v1/predict
- **Method:** POST
- **Content-Type:** application/json
- **Authorization:** Bearer <JWT>

Sample JSON Request

[2 items]

```
customer_id: cust-a1b2c3d4  
behavioral_score: 75  
transactional_value_k: 120  
support_ticket_volume_30d: 1  
engagement_score: 0.92  
contract_days_remaining: 360
```

```
customer_id: cust-e5f6g7h8  
behavioral_score: 30.5  
transactional_value_k: 550  
support_ticket_volume_30d: 9  
engagement_score: 0.25  
contract_days_remaining: 45
```

Sample JSON Response (HTTP 200 OK)

[2 items]

```
customer_id: cust-a1b2c3d4
churn_probability: 0.0512
churn_risk_level: LOW
action_code: ACTION-LOW
top_churn_feature: contract_days_remaining
```

```
customer_id: cust-e5f6g7h8
churn_probability: 0.8255
churn_risk_level: HIGH
action_code: ACTION-HIGH
top_churn_feature: support_ticket_volume_30d
```

Simple SQL DDL for Predictions Log Table

SQL

```
CREATE TABLE predictions_log (
    prediction_id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    customer_id VARCHAR(128) NOT NULL,
    prediction_timestamp TIMESTAMP WITH TIME ZONE DEFAULT NOW(),
    churn_probability REAL NOT NULL,
    churn_risk_level VARCHAR(10) NOT NULL,
    model_version VARCHAR(32) NOT NULL,
    inference_features JSONB NOT NULL,
    webhook_status VARCHAR(50) DEFAULT 'PENDING'
);
```

```
CREATE INDEX idx_customer_time ON predictions_log (customer_id, prediction_timestamp DESC);
```

RACI Matrix for Deployment Responsibilities

Activity	DevOps Engineer	MLOps Engineer	Enterprise Architect	CEO (Alex Rojas)
Infrastructure Provisioning (EKS/Helm)	Responsible/Accountable	Consulted	Informed	Informed
Model Deployment & Validation	Consulted	Responsible/Accountable	Informed	Informed
CI/CD Pipeline Maintenance	Responsible	Responsible	Informed	Informed
Security & Compliance Review	Informed	Informed	Accountable	Consulted
Strategic Business Alignment	Informed	Informed	Responsible	Accountable